

Assignment No 3

Aim:- Build the image classification model by dividing the model into 4 stages

- a] Loading image data
- b] Defining model's architecture
- c] Training the model
- d] Estimate model's performance

Objectives:- To learn about CNN & how to develop a CNN for image recognition.

Infrastructure: Computer / Laptop

Software used:- Jupyter Notebook / Google Collab

Theory:-

CNN

Convolution networks, also known as convolution neural networks or CNN's are a specialized kind of neural network for processing data that has a known, grid-like topology. The name "convolution neural network" indicates that the network employs a mathematical operation called convolution. Convolution is a special kind of linear operation. Convolution networks are simply neural networks that use convolution in place of general matrix multiplication in at least one of their layers.

- These are two main parts to a CNN architecture.
- A convolution tool that separates and identifies the various features of that image for analysis in a process called as Feature Extraction.
 - The network of feature extraction consists of multiple pairs of convolutional or pooling layers.
 - A fully connected layer that utilizes the output from the convolutional process & predicts the class of image based on feature extracted in previous steps.
 - This CNN model of feature extraction aims to reduce the number of features present in a dataset. It creates new features which summarise the existing features contained in an original set of features.

Convolution Layers

- There are three types of layers that make up a CNN: which are convolution layers, pooling layers & fully connected layers.
- In addition to these three layers there are two important parameters which are dropout layers & activation function.

1] Convolution layer :-

This is the first layer that is used to extract various features from input images. In this layer, mathematical operations are performed between an input image & a filter. By sliding over the input image, the dot product is taken between the filters and the input image with respect to size of filter.

• The output is termed as the Feature map which gives us information about image such as corners & edges.

2] Pooling Layers

The primary aim of this layer is to decrease the size of the convolved feature map to reduce the computational costs. This is performed by decreasing the connections between layers & independently operates on each feature map. In Max Pooling, the largest element is taken from feature map. Average pooling calculates the average of elements in predefined sized image.

3] Fully Connected Layer

The fully connected layer consists of weights & biases along with neuron & is used to connect the neurons between two different layers. These layers are usually placed before the output layer & form the last few layers of an CNN architecture.

4] Dropout :-

Usually, when all features are connected to FC layer, it can cause overfitting in the training dataset. Overfitting occurs when a particular model works so well on training data causing a negative impact in model's performance when used on a new data.

5.] Activation Function

They are used to learn & approximate any of continuous & complex relationship between variables of network. In simple words, it decides which information of the model should fire in forward direction & which ones should not at end of network.

Implementation

- 1.] Load necessary libraries
- 2.] Import dataset from respective library
- 3.] Design neural network architecture & mention number of layers, etc.
- 4.] Train the model with imported dataset
- 5.] Evaluate performance of model.

Conclusion :-

We learnt how to build & train a CNN to identify images.


```
In [1]: from tensorflow.keras.models import Sequential
        from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dense, Flatten, Dropout
        from tensorflow.keras.optimizers import SGD
        import numpy as np
        import random
        import matplotlib.pyplot as plt
        import tensorflow as tf
```

```
In [2]: import os
        os.environ['CUDA_VISIBLE_DEVICES'] = '-1'
```

Load Dataset

Dataset available on <https://bit.ly/ImgClsKeras> (<https://bit.ly/ImgClsKeras>)

```
In [3]: x_train = np.loadtxt('input.csv', delimiter = ',')
        y_train = np.loadtxt('labels.csv', delimiter = ',')

        x_test = np.loadtxt('input_test.csv', delimiter = ',')
        y_test = np.loadtxt('labels_test.csv', delimiter = ',')
```

```
In [4]: x_train = x_train.reshape(len(x_train), 100, 100, 3)
        y_train = y_train.reshape(len(y_train), 1)

        x_test = x_test.reshape(len(x_test), 100, 100, 3)
        y_test = y_test.reshape(len(y_test), 1)

        x_train = x_train/255.0
        x_test = x_test/255.0
```

```
In [5]: print("Shape of X_train:", x_train.shape)
        print("Shape of Y_train:", y_train.shape)
        print("Shape of X_train:", x_test.shape)
        print("Shape of X_train:", y_test.shape)

        Shape of X_train: (2000, 100, 100, 3)
        Shape of Y_train: (2000, 1)
        Shape of X_train: (400, 100, 100, 3)
        Shape of X_train: (400, 1)
```

```
In [6]: idx = random.randint(0,len(x_train))
plt.imshow(x_train[idx,:])
plt.show()
```



Model Building

```
In [8]: model = Sequential([
    Conv2D(256,(3,3),activation = 'relu',input_shape=(100,100,3)),
    BatchNormalization(),
    MaxPooling2D((4,4)),
    Conv2D(128,(3,3),activation = 'relu'),
    BatchNormalization(),
    MaxPooling2D((2,2)),
    Conv2D(64,(3,3),activation = 'relu'),
    BatchNormalization(),
    MaxPooling2D((2,2)),

    Flatten(),
    Dense(128,activation='relu'),
    Dropout(0.4),
    Dense(1,activation='sigmoid')
])
```

```
In [9]: opt = SGD(momentum=0.9)
model.compile(optimizer=opt, loss='binary_crossentropy', metrics = ['accuracy'])
```

```
In [10]: model.fit(x_train,y_train,epochs=30,steps_per_epoch = 20,validation_data=(x_test,
```

```
Epoch 1/30
```

```
20/20 [=====] - 61s 3s/step - loss: 0.8490 - accuracy:  
0.5650 - val_loss: 0.6948 - val_accuracy: 0.5250
```

```
Epoch 2/30
```

```
20/20 [=====] - 70s 4s/step - loss: 0.6606 - accuracy:  
0.6240 - val_loss: 0.6801 - val_accuracy: 0.5975
```

```
Epoch 3/30
```

```
20/20 [=====] - 68s 3s/step - loss: 0.6057 - accuracy:  
0.6830 - val_loss: 0.6797 - val_accuracy: 0.5625
```

```
Epoch 4/30
```

```
20/20 [=====] - 64s 3s/step - loss: 0.5674 - accuracy:  
0.6990 - val_loss: 0.7254 - val_accuracy: 0.5050
```

```
Epoch 5/30
```

```
20/20 [=====] - 64s 3s/step - loss: 0.5089 - accuracy:  
0.7605 - val_loss: 0.7999 - val_accuracy: 0.5050
```

```
Epoch 6/30
```

```
20/20 [=====] - 62s 3s/step - loss: 0.4879 - accuracy:  
0.7720 - val_loss: 0.8164 - val_accuracy: 0.4950
```

```
Epoch 7/30
```

```
20/20 [=====] - 61s 3s/step - loss: 0.4430 - accuracy:  
0.7960 - val_loss: 0.8590 - val_accuracy: 0.5125
```

```
Epoch 8/30
```

```
20/20 [=====] - 60s 3s/step - loss: 0.4141 - accuracy:  
0.8130 - val_loss: 0.8978 - val_accuracy: 0.5050
```

```
Epoch 9/30
```

```
20/20 [=====] - 65s 3s/step - loss: 0.3859 - accuracy:  
0.8270 - val_loss: 1.0060 - val_accuracy: 0.5025
```

```
Epoch 10/30
```

```
20/20 [=====] - 59s 3s/step - loss: 0.3405 - accuracy:  
0.8470 - val_loss: 1.0907 - val_accuracy: 0.5200
```

```
Epoch 11/30
```

```
20/20 [=====] - 59s 3s/step - loss: 0.3248 - accuracy:  
0.8555 - val_loss: 0.9742 - val_accuracy: 0.5300
```

```
Epoch 12/30
```

```
20/20 [=====] - 59s 3s/step - loss: 0.2908 - accuracy:  
0.8820 - val_loss: 0.8911 - val_accuracy: 0.5650
```

```
Epoch 13/30
```

```
20/20 [=====] - 59s 3s/step - loss: 0.2615 - accuracy:  
0.8920 - val_loss: 0.6965 - val_accuracy: 0.6225
```

```
Epoch 14/30
```

```
20/20 [=====] - 59s 3s/step - loss: 0.2234 - accuracy:  
0.9145 - val_loss: 0.7834 - val_accuracy: 0.6000
```

```
Epoch 15/30
```

```
20/20 [=====] - 60s 3s/step - loss: 0.1961 - accuracy:  
0.9245 - val_loss: 1.3311 - val_accuracy: 0.5475
```

```
Epoch 16/30
```

```
20/20 [=====] - 60s 3s/step - loss: 0.1657 - accuracy:  
0.9470 - val_loss: 1.0894 - val_accuracy: 0.5900
```

```
Epoch 17/30
```

```
20/20 [=====] - 60s 3s/step - loss: 0.1325 - accuracy:  
0.9590 - val_loss: 1.3612 - val_accuracy: 0.5525
```

```
Epoch 18/30
```

```
20/20 [=====] - 60s 3s/step - loss: 0.1255 - accuracy:  
0.9565 - val_loss: 1.0850 - val_accuracy: 0.6050
```

```
Epoch 19/30
```



```

20/20 [=====] - 60s 3s/step - loss: 0.1015 - accuracy:
0.9620 - val_loss: 1.1502 - val_accuracy: 0.5850
Epoch 20/30
20/20 [=====] - 63s 3s/step - loss: 0.0923 - accuracy:
0.9660 - val_loss: 1.0095 - val_accuracy: 0.6275
Epoch 21/30
20/20 [=====] - 61s 3s/step - loss: 0.0725 - accuracy:
0.9770 - val_loss: 1.1029 - val_accuracy: 0.6150
Epoch 22/30
20/20 [=====] - 60s 3s/step - loss: 0.0776 - accuracy:
0.9725 - val_loss: 0.7184 - val_accuracy: 0.7275
Epoch 23/30
20/20 [=====] - 61s 3s/step - loss: 0.1011 - accuracy:
0.9610 - val_loss: 0.9714 - val_accuracy: 0.6525
Epoch 24/30
20/20 [=====] - 61s 3s/step - loss: 0.0831 - accuracy:
0.9705 - val_loss: 0.9042 - val_accuracy: 0.6775
Epoch 25/30
20/20 [=====] - 61s 3s/step - loss: 0.0670 - accuracy:
0.9760 - val_loss: 1.0374 - val_accuracy: 0.6825
Epoch 26/30
20/20 [=====] - 61s 3s/step - loss: 0.0483 - accuracy:
0.9835 - val_loss: 0.8355 - val_accuracy: 0.7250
Epoch 27/30
20/20 [=====] - 61s 3s/step - loss: 0.0448 - accuracy:
0.9860 - val_loss: 1.0373 - val_accuracy: 0.7000
Epoch 28/30
20/20 [=====] - 61s 3s/step - loss: 0.0401 - accuracy:
0.9870 - val_loss: 0.8931 - val_accuracy: 0.7225
Epoch 29/30
20/20 [=====] - 61s 3s/step - loss: 0.0376 - accuracy:
0.9900 - val_loss: 1.1096 - val_accuracy: 0.7050
Epoch 30/30
20/20 [=====] - 61s 3s/step - loss: 0.0272 - accuracy:
0.9940 - val_loss: 1.0204 - val_accuracy: 0.7100

```

```
Out[10]: <keras.callbacks.History at 0x1c39aeacee0>
```

```
In [11]: model.evaluate(x_test,y_test)
```

```

13/13 [=====] - 2s 164ms/step - loss: 1.0204 - accurac
y: 0.7100

```

```
Out[11]: [1.0204460620880127, 0.7099999785423279]
```

Making Predictions


```


In [14]: idx2 = random.randint(0, len(y_test))
plt.imshow(x_test[idx2,:])
plt.show()

y_pred = model.predict(x_test[idx2,:].reshape(1,100,100,3))
y_pred = y_pred > 0.5

if(y_pred==0):
    pred = 'dog'
else:
    pred = 'cat'

print("Our model says it is a",pred)

```



```

1/1 [=====] - 0s 28ms/step
Our model says it is a dog

```

```
In [15]: score = model.evaluate(x_test, y_test, verbose = 0 )
print("Test Score: ", score[0])
print("Test accuracy: ", score[1])
```

```
Test Score: 1.0204460620880127
Test accuracy: 0.7099999785423279
```

```
In [16]: model.summary()
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 98, 98, 256)	7168
batch_normalization (Batch Normalization)	(None, 98, 98, 256)	1024
max_pooling2d (MaxPooling2D)	(None, 24, 24, 256)	0
conv2d_1 (Conv2D)	(None, 22, 22, 128)	295040
batch_normalization_1 (Batch Normalization)	(None, 22, 22, 128)	512
max_pooling2d_1 (MaxPooling2D)	(None, 11, 11, 128)	0
conv2d_2 (Conv2D)	(None, 9, 9, 64)	73792
batch_normalization_2 (Batch Normalization)	(None, 9, 9, 64)	256
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 64)	0
flatten (Flatten)	(None, 1024)	0
dense (Dense)	(None, 128)	131200
dropout (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 1)	129

```
=====
Total params: 509,121
Trainable params: 508,225
Non-trainable params: 896
```



```

In [17]: val = model.fit(x_train,y_train, epochs=5,validation_data=(x_test,y_test),batch_

Epoch 1/5
10/10 [=====] - 61s 6s/step - loss: 0.0246 - accuracy:
0.9925 - val_loss: 0.9134 - val_accuracy: 0.7250
Epoch 2/5
10/10 [=====] - 59s 6s/step - loss: 0.0173 - accuracy:
0.9970 - val_loss: 0.8590 - val_accuracy: 0.7275
Epoch 3/5
10/10 [=====] - 58s 6s/step - loss: 0.0139 - accuracy:
0.9970 - val_loss: 0.8365 - val_accuracy: 0.7425
Epoch 4/5
10/10 [=====] - 59s 6s/step - loss: 0.0101 - accuracy:
0.9995 - val_loss: 0.8777 - val_accuracy: 0.7425
Epoch 5/5
10/10 [=====] - 59s 6s/step - loss: 0.0116 - accuracy:
0.9975 - val_loss: 0.8869 - val_accuracy: 0.7650

```

```

In [18]: plt.title("Model Accuracy")
plt.ylabel("Accuracy")
plt.xlabel('epoch')
plt.plot(val.history['accuracy'])
plt.plot(val.history['val_accuracy'])
plt.legend(['train','val'])
plt.show()

```

